

ASSIGNMENT 6

- Name : Achal Rajesh Mate
- Roll No : 2203541
- Enroll No : MITU20BTCSD001
- Branch : CSE
- Class : TY CSE Is - 3
- Guided By : Prof Nagesh Jadhav Sir

Title:- Implement K means algorithm on dataset

Objectives:-

1. To learn unsupervised learning
2. To implement K means algorithm

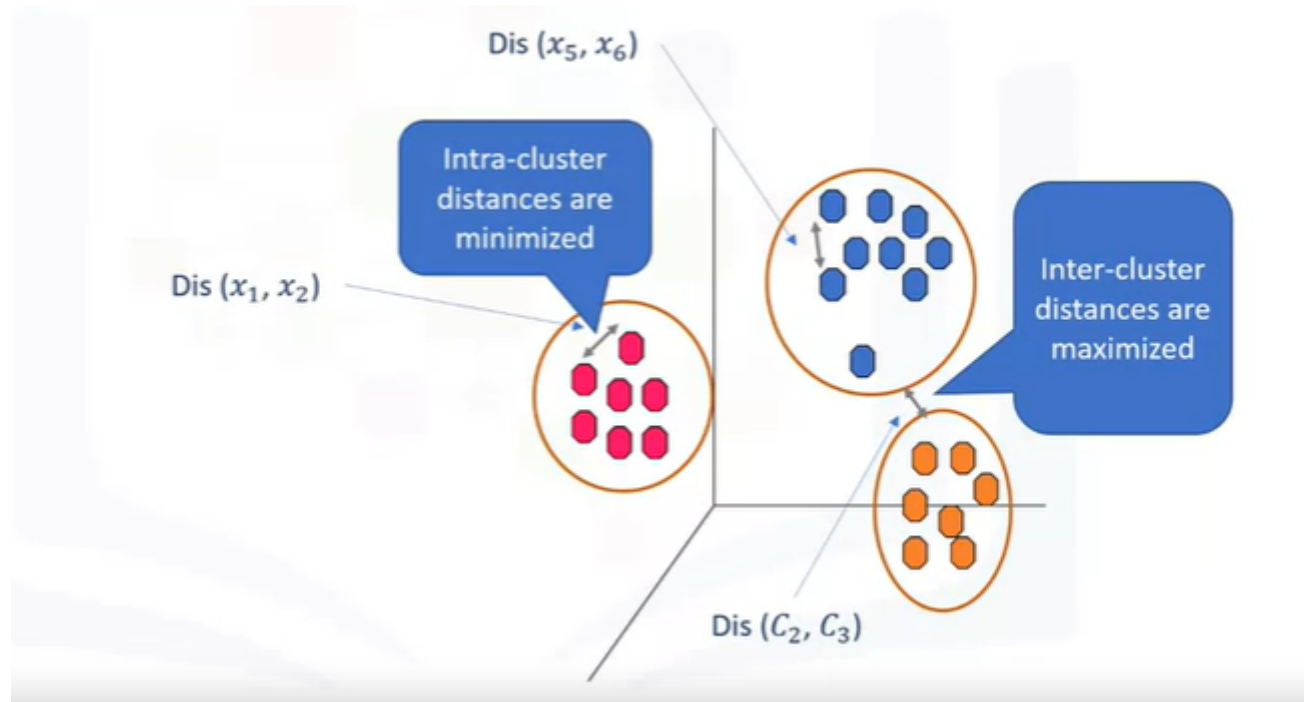
Theory:

K-Means Algorithm?

- K-Means Clustering is an Unsupervised Learning algorithm, which groups the unlabeled dataset into different clusters. Here K defines the number of pre-defined clusters that need to be created in the process, as if K=2, there will be two clusters, and for K=3, there will be three clusters, and so on.
- It allows us to cluster the data into different groups
- It is a centroid-based algorithm, where each cluster is associated with a centroid.
- The main aim of this algorithm is to minimize the sum of distances between the data point and their corresponding clusters.

$$\text{Dis}(x_1, x_2) = \sqrt{\sum_{i=0}^n (x_{1i} - x_{2i})^2}$$

- Distance of samples from each other is used to shape the cluster




- The algorithm takes the unlabeled dataset as input, divides the dataset into k-number of clusters, and repeats the process until it does not find the best clusters. The value of k should be predetermined in this algorithm.

The k-means clustering algorithm mainly performs two tasks:

- Determines the best value for K center points or centroids by an iterative process.
- Assigns each data point to its closest k-center. Those data points which are near to the particular k-center, create a cluster.
- Hence each cluster has datapoints with some commonalities, and it is away from other clusters.

The below diagram explains the working of the K-means Clustering Algorithm:



Steps to implement K-Means Algorithm

- Step-1: Select the number K to decide the number of clusters.
- Step-2: Select random K points or centroids. (It can be other from the input dataset).
- Step-3: Assign each data point to their closest centroid, which will form the predefined K clusters.
- Step-4: Calculate the variance and place a new centroid of each cluster.
- Step-5: Repeat the third steps, which means reassign each datapoint to the new closest centroid of each cluster.
- Step-6: If any reassignment occurs, then go to step-4 else go to FINISH.
- Step-7: The model is ready.

RULE FOR CHOOSING VALUE OF K

- Never go with $k=2$ as the option because it means you divide the complete data into two halves and it's not useful for any business case.
- If you have option to choose between two values of K, always go with a lesser value.
- Since we will be taking business decisions based on the cluster result, it's always a good idea to go with a lower value of K so that it's easy to take and implement business decisions.
- Silhouette: That value of k for which the score is maximum

- Elbow, you look at the elbow of the curve

Problem Statement : Customer Segmentation

Prolem Context :

Customer segmentation is the practice of partitioning a customer base into groups of individuals that have similar characteristics. It is a significant strategy as a business can target these specific groups of customers and effectively allocate marketing resources. For example, one group might contain customers who are high-profit and low-risk, that is, more likely to purchase products, or subscribe for a service. A business task is to retain those customers. Another group might include customers from non-profit organizations and so on.

Dataset Link :

https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-ML0101EN-SkillsNetwork/labs/Module%204/data/Cust_Segmentation.csv

Dataset Information

This Dataset Contains 10 features with 850 instances Attributes are

1. Customer Id,
2. Age
3. Edu
4. Years Employed
5. Income
6. Card Debt
7. Other Debt
8. Defaulted
9. Address 10. DebtIncomeRatio

Import All Necessary Files

```
In [1]: import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

```
In [31]: from sklearn.cluster import KMeans
from sklearn.preprocessing import LabelEncoder, StandardScaler
from scipy import stats
from sklearn.metrics import silhouette_score, accuracy_score, confusion_matrix, classification_report
from sklearn import metrics
```

Read the Dataset

```
In [3]: df = pd.read_csv("Cust_Segmentation.csv")
df.head()
```

```
Out[3]:
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	Defaulted	Address	DebtIncomeRatio
0	1	41	2	6	19	0.124	1.073	0.0	NBA001	6.3
1	2	47	1	26	100	4.582	8.218	0.0	NBA021	12.8
2	3	33	2	10	57	6.111	5.802	1.0	NBA013	20.9
3	4	29	2	4	19	0.681	0.516	0.0	NBA009	6.3
4	5	47	1	31	253	9.308	8.908	0.0	NBA008	7.2

View Dimension of dataste

```
In [4]: df.shape
```

```
Out[4]: (850, 10)
```

This dataset contains 850 instance or rows and 10 columns

Columns in dataset

```
In [5]: df.columns
```

```
Out[5]: Index(['Customer Id', 'Age', 'Edu', 'Years Employed', 'Income', 'Card Debt',  
             'Other Debt', 'Defaulted', 'Address', 'DebtIncomeRatio'],  
            dtype='object')
```

Concise Summary

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 850 entries, 0 to 849  
Data columns (total 10 columns):  
#   Column                Non-Null Count  Dtype    
---  ---                     -  
0   Customer Id           850 non-null    int64    
1   Age                   850 non-null    int64    
2   Edu                   850 non-null    int64    
3   Years Employed        850 non-null    int64    
4   Income                850 non-null    int64    
5   Card Debt             850 non-null    float64  
6   Other Debt            850 non-null    float64  
7   Defaulted             700 non-null    float64  
8   Address               850 non-null    object    
9   DebtIncomeRatio       850 non-null    float64  
dtypes: float64(4), int64(5), object(1)  
memory usage: 66.5+ KB
```

from above information we can say that Address feature is in object datatype and defaulted feature contains null values ,Lets check them

```
In [7]: df.isnull().sum()
```

```
Out[7]: Customer Id      0  
Age                    0  
Edu                    0  
Years Employed         0  
Income                 0  
Card Debt              0  
Other Debt             0
```

```
Defaulted      150
Address         0
DebtIncomeRatio 0
dtype: int64
```

defaulted feature contains 150 NULL values

Drop the Column That Contains Null Values and Categorical Values

The k-means algorithm isn't directly applicable to categorical variables because the Euclidean distance function isn't really meaningful for discrete variables.

```
In [8]: df.Address.value_counts()
```

```
Out[8]: NBA001      71
NBA002      71
NBA000      60
NBA004      58
NBA003      55
NBA006      50
NBA008      49
NBA009      45
NBA005      43
NBA007      41
NBA010      37
NBA011      36
NBA012      28
NBA014      24
NBA016      22
NBA013      22
NBA017      20
NBA015      18
NBA019      16
NBA018      14
NBA023      11
NBA021      10
NBA026      10
NBA022       9
NBA025       9
NBA020       8
NBA024       4
NBA027       4
NBA031       2
NBA034       1
NBA029       1
```

```
NBA030      1  
Name: Address, dtype: int64
```

```
In [9]: df.drop(['Customer Id', 'Defaulted', 'Address'], axis=1, inplace=True)
```


```
In [10]: df.columns
```

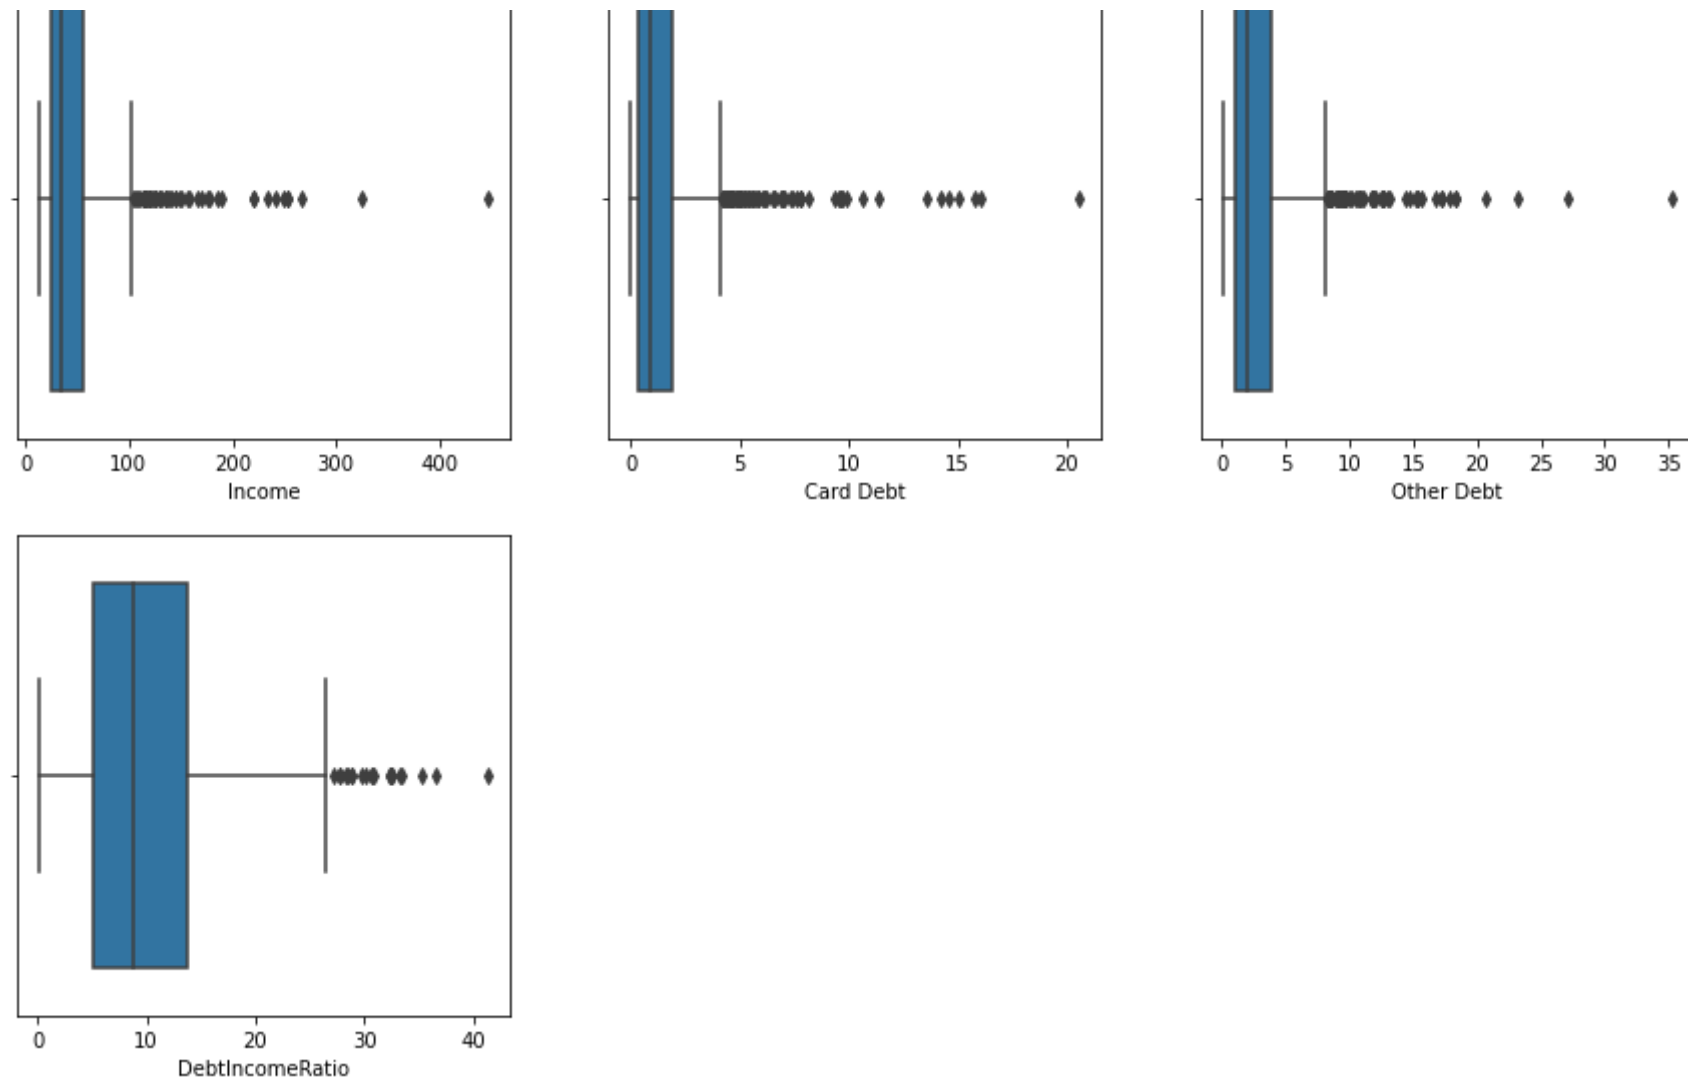
```
Out[10]: Index(['Age', 'Edu', 'Years Employed', 'Income', 'Card Debt', 'Other Debt',  
              'DebtIncomeRatio'],  
              dtype='object')
```

Check the outlier by plotting box plot

```
In [11]: f = df.columns[:]
```

```
In [12]: plt.figure(figsize=(15,15))  
for col in enumerate(f):  
    plt.subplot(3,3,col[0] + 1)  
    sns.boxplot(data=df, x=col[1])
```





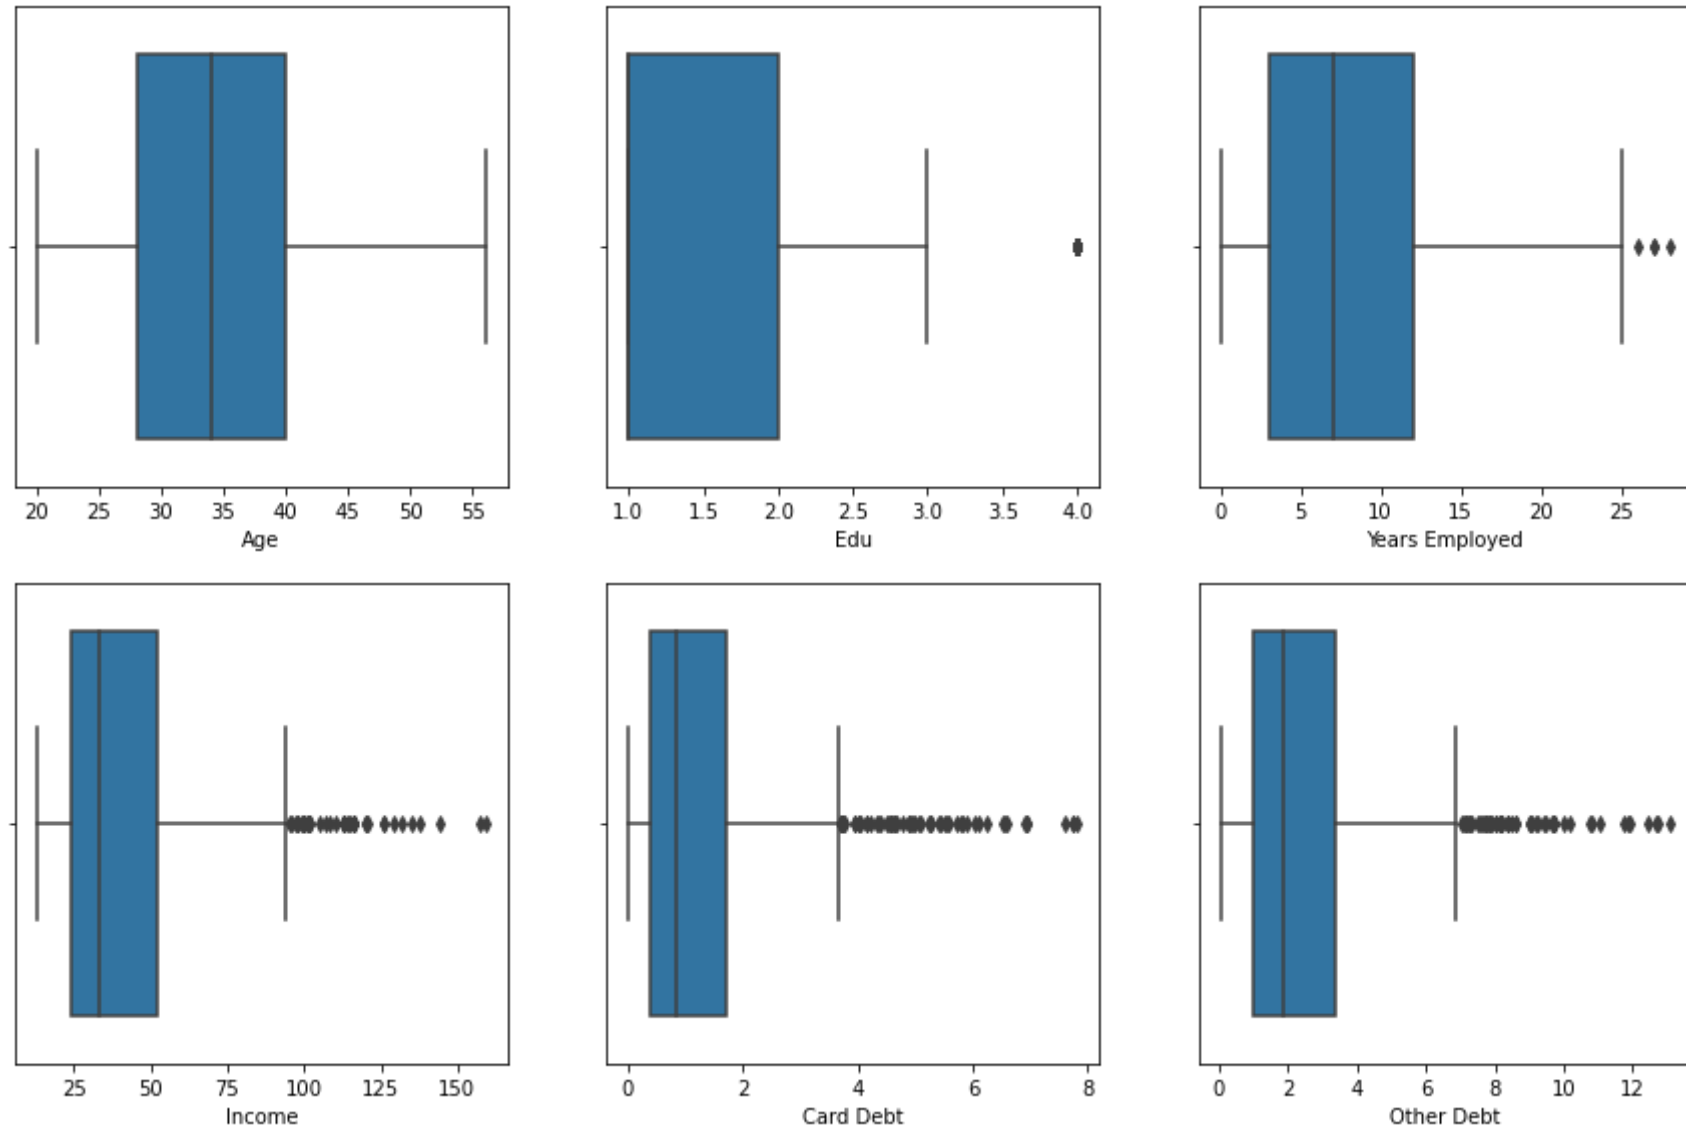
Drop The Outlier Using Z score


```
In [13]: z = np.abs(stats.zscore(df))
df = df[(z<3).all(axis=1)]
df.shape
```

Out[13]: (802, 7)

48 rows are drop when outliers are removed using Z-score

```
In [14]: plt.figure(figsize=(15,15))  
for col in enumerate(f):  
    plt.subplot(3,3,col[0] + 1)  
    sns.boxplot(data=df, x=col[1])
```





Data Normalizing over the standard deviation

```
In [16]: df_scalar = StandardScaler().fit_transform(df)
```

```
In [17]: df_scalar = pd.DataFrame(df_scalar ,columns=df.columns)
```

```
In [18]: df_scalar
```

```
Out[18]:
```

	Age	Edu	Years Employed	Income	Card Debt	Other Debt	DebtIncomeRatio
0	0.816545	0.373949	-0.333197	-0.906959	-0.850017	-0.659345	-0.568286
1	1.579796	-0.766382	2.872850	2.367348	2.375862	2.425850	0.520179
2	-0.201123	0.373949	0.308012	0.629135	3.482271	1.382627	1.876575
3	-0.709957	0.373949	-0.653802	-0.906959	-0.446963	-0.899856	-0.568286
4	0.689337	-0.766382	2.391943	1.599300	-0.217576	2.258744	0.202012
...
797	0.816545	-0.766382	-0.172895	0.063206	-0.437556	-0.605370	-0.886453

	Age	Edu	Years Employed	Income	Card Debt	Other Debt	DebtIncomeRatio
798	-0.964374	-0.766382	-0.493499	-0.623994	-0.543204	-0.595871	-0.484558
799	-0.837165	0.373949	-0.172895	-0.300606	-0.679967	-0.250001	-0.451067
800	-0.328331	-0.766382	0.628617	-0.543147	-0.855806	-0.822133	-1.137638
801	2.215839	-0.766382	1.269827	0.912100	0.410522	0.448216	-0.183137

802 rows × 7 columns

Model Building

In [19]:

```
clusterNum = 3
k_means = KMeans(init = "k-means++", n_clusters = clusterNum, n_init = 12)
k_means.fit(df_scalar)
labels = k_means.labels_
print(labels)
```

```
[1 2 2 1 2 0 0 1 0 0 1 1 0 1 1 1 1 0 0 1 0 2 0 0 1 0 0 1 0 2 1 1 1 1 1 1 1
2 1 2 2 2 2 1 1 1 0 0 1 0 2 1 1 1 0 0 2 0 2 2 1 1 0 1 1 1 1 0 2 1 1 1 0 1
0 2 2 1 2 1 1 1 1 1 0 1 1 1 1 0 1 0 1 1 2 2 0 0 1 2 0 1 0 0 0 1 1 1 1 0 1
0 2 1 1 1 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 1 0 1 2 1 1 1 2 1 0 2 1 1 0 2 1 1
0 2 0 0 1 2 1 2 0 0 1 1 1 1 1 0 2 1 1 1 1 2 0 1 1 0 2 0 1 1 2 1 0 0 0 1 2
0 1 0 1 1 0 1 2 1 1 1 0 2 1 0 1 1 0 0 0 1 1 2 2 0 0 1 0 1 1 1 1 2 1 2 2 0
1 2 1 2 1 0 1 2 1 1 0 1 1 1 0 2 0 1 1 1 0 2 0 1 1 1 1 2 2 1 2 0 1 2 1
0 1 1 0 0 1 1 1 2 1 1 1 0 1 2 1 2 1 0 2 1 0 1 2 1 1 1 1 2 0 0 1 1 1 0 2
1 1 1 1 1 1 0 1 1 1 0 1 1 2 1 1 0 1 1 0 0 0 2 1 0 0 0 1 2 1 0 0 0 0 1 2 0
1 1 1 1 1 0 1 1 0 2 1 1 1 1 1 1 0 1 0 2 0 1 1 0 0 1 1 2 0 1 1 1 1 1 0 2 0
0 1 1 1 1 2 0 1 0 1 0 1 1 1 1 1 1 1 2 0 2 0 0 1 0 2 1 1 0 2 1 2 1 0 2 1 1
2 1 0 1 1 2 2 1 1 2 1 1 1 2 0 0 0 1 0 2 0 0 0 0 1 1 1 1 1 1 1 2 1 1 1 1 1
0 1 0 2 2 1 0 0 0 1 0 1 1 1 1 1 0 1 1 1 2 2 1 1 1 2 0 2 0 1 1 0 1 1 1 2 0
1 1 0 2 0 1 1 1 1 0 0 1 1 1 2 2 1 1 1 0 1 1 0 1 1 1 2 1 1 1 1 1 2 0 0 1 0
0 0 0 1 0 1 2 1 0 1 0 0 1 1 0 0 1 2 1 0 2 1 1 1 1 1 0 1 1 1 1 0 1 0 0 0 2
1 1 2 1 2 0 2 1 1 1 0 0 0 0 1 2 1 2 1 1 1 1 1 1 1 2 1 2 1 1 1 1 1 1 1 2
1 1 0 0 1 0 0 2 1 1 2 1 0 1 1 1 0 2 0 1 0 1 1 1 2 2 1 1 1 1 1 1 0 2 1 0 0
1 1 1 0 0 1 1 0 1 0 1 1 2 1 2 0 1 0 1 0 1 1 1 0 1 2 1 1 2 0 1 1 1 1 1 2
1 1 2 1 2 1 2 0 2 1 0 0 2 1 1 1 1 1 1 1 1 2 2 1 0 0 1 1 1 0 2 0 1 2 1 0
0 2 1 1 1 1 1 1 1 1 1 0 2 2 2 1 0 1 1 0 0 1 0 1 1 1 0 2 0 1 0 2 1 1 0 0
2 1 0 1 1 1 0 2 2 1 1 1 1 1 1 1 1 1 1 0 1 2 1 1 1 1 1 1 1 1 1 1 1 2 1
1 1 2 1 1 1 2 1 2 1 1 1 1 1 1 1 0 1 0 1 0 1 1 1 0]
```


Let's run K-Means with different value of K to check SILHOUETTE SCORE

```
In [20]: sil = []  
         for k in range(2,11):  
             kmean = KMeans(n_clusters=k, random_state=0,init = "k-means++", n_init = 12).fit(df_scalar)  
             sil.append([k,silhouette_score(df_scalar,kmean.labels_)])
```

```
In [21]: sil
```

```
Out[21]: [[2, 0.32447376073926576],  
          [3, 0.23747510415997428],  
          [4, 0.2186339487233202],  
          [5, 0.22942591581160304],  
          [6, 0.20742263869379915],  
          [7, 0.21293537439736465],  
          [8, 0.20401959301325623],  
          [9, 0.20499215165510967],  
          [10, 0.18406829643929346]]
```

```
In [22]: sil = pd.DataFrame(sil)  
         plt.figure(figsize=(10,10))  
         plt.plot(sil[0], sil[1])  
         plt.show()
```




Choosing Value of K

K value with 2 has maximum Silhouette Score is 0.32 but choosing $k=2$ is not good choice because it means you divide the complete data into two halves and it's not useful for any business case.
Second Largest Silhouette Score is 0.23 with $k=3$, so we can go with $k=3$

Let's run K-Means with different value of K using Elbow Method

```
In [23]: ssd = []
for k in range(2,11):
    kmean_elbow = KMeans(n_clusters=k).fit(df_scalar)
    ssd.append([k, kmean_elbow.inertia_])
ssd = pd.DataFrame(ssd)
plt.plot(ssd[0], ssd[1])
plt.show()
```



In Elbow method we look at the elbow of the curve here K= 3 form elbow of curve so we select k = 3

```
In [50]: kmean = KMeans(n_clusters=3, random_state=0,init = "k-means++", n_init = 12)
kmean.fit(df_scalar)
```

```
Out[50]: KMeans(n_clusters=3, n_init=12, random_state=0)
```

```
In [51]: labels = k_means.labels_
print(labels)
```

```
[1 2 2 1 2 0 0 1 0 0 1 1 0 1 1 1 1 0 0 1 0 2 0 0 1 0 0 1 0 2 1 1 1 1 1 1 1
 2 1 2 2 2 2 1 1 1 0 0 1 0 2 1 1 1 0 0 2 0 2 2 1 1 0 1 1 1 1 0 2 1 1 1 0 1
 0 2 2 1 2 1 1 1 1 1 0 1 1 1 1 0 1 0 1 1 2 2 0 0 1 2 0 1 0 0 0 1 1 1 1 0 1
 0 2 1 1 1 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 1 0 1 2 1 1 1 2 1 0 2 1 1 0 2 1 1]
```

```

0 2 0 0 1 2 1 2 0 0 1 1 1 1 1 0 2 1 1 1 2 0 1 1 0 2 0 1 1 2 1 0 0 0 1 2
0 1 0 1 1 0 1 2 1 1 1 0 2 1 0 1 1 0 0 0 1 1 2 2 0 0 1 0 1 1 1 1 2 1 2 2 0
1 2 1 2 1 0 1 2 1 1 0 1 1 1 0 2 0 1 1 1 0 1 0 2 0 1 1 1 1 2 2 1 2 0 1 2 1
0 1 1 0 0 1 1 1 2 1 1 1 0 1 2 1 2 1 0 2 1 0 1 2 1 1 1 1 1 2 0 0 1 1 1 0 2
1 1 1 1 1 1 0 1 1 1 0 1 1 2 1 1 0 1 1 0 0 0 2 1 0 0 0 1 2 1 0 0 0 0 1 2 0
1 1 1 1 1 0 1 1 0 2 1 1 1 1 1 1 0 1 0 2 0 1 1 0 0 1 1 2 0 1 1 1 1 1 0 2 0
0 1 1 1 1 2 0 1 0 1 0 1 1 1 1 1 1 1 2 0 2 0 0 1 0 2 1 1 0 2 1 2 1 0 2 1 1
2 1 0 1 1 2 2 1 1 2 1 1 1 2 0 0 0 1 0 2 0 0 0 0 1 1 1 1 1 1 2 1 1 1 1 1
0 1 0 2 2 1 0 0 0 1 0 1 1 1 1 1 0 1 1 1 2 2 1 1 1 2 0 2 0 1 1 0 1 1 1 2 0
1 1 0 2 0 1 1 1 1 0 0 1 1 1 2 2 1 1 1 0 1 1 0 1 1 1 2 1 1 1 1 2 0 0 1 0
0 0 0 1 0 1 2 1 0 1 0 0 1 1 0 0 1 2 1 0 2 1 1 1 1 1 0 1 1 1 0 1 0 0 0 2
1 1 2 1 2 0 2 1 1 1 0 0 0 0 1 2 1 2 1 1 1 1 1 1 1 2 1 2 1 1 1 1 1 1 1 2
1 1 0 0 1 0 0 2 1 1 2 1 0 1 1 1 0 2 0 1 0 1 1 1 2 2 1 1 1 1 1 0 2 1 0 0
1 1 1 0 0 1 1 0 1 0 1 1 2 1 2 0 1 0 1 0 1 1 1 0 1 2 1 1 2 0 1 1 1 1 1 2
1 1 2 1 2 1 2 0 2 1 0 0 2 1 1 1 1 1 1 1 1 2 2 1 0 0 1 1 1 0 2 0 1 2 1 0
0 2 1 1 1 1 1 1 1 1 1 0 2 2 2 1 0 1 1 0 0 1 0 1 1 1 0 2 0 1 0 2 1 1 0 0
2 1 0 1 1 1 0 2 2 1 1 1 1 1 1 1 1 1 1 0 1 2 1 1 1 1 1 1 1 1 1 1 1 2 1
1 1 2 1 1 1 2 1 2 1 1 1 1 1 1 1 0 1 0 1 0 1 1 1 0]

```

Assign the labels to each row in the dataframe.

In [52]:

```

df_scalar["Label"] = labels
df_scalar.head(5)

```

Out[52]:

	Age	Edu	Years Employed	Income	Card Debt	Other Debt	DebtIncomeRatio	Label
0	0.816545	0.373949	-0.333197	-0.906959	-0.850017	-0.659345	-0.568286	1
1	1.579796	-0.766382	2.872850	2.367348	2.375862	2.425850	0.520179	2
2	-0.201123	0.373949	0.308012	0.629135	3.482271	1.382627	1.876575	2
3	-0.709957	0.373949	-0.653802	-0.906959	-0.446963	-0.899856	-0.568286	1
4	0.689337	-0.766382	2.391943	1.599300	-0.217576	2.258744	0.202012	2

In [53]:

```

df_scalar.Label.value_counts()

```

Out[53]:

```

1    453
0    212
2    137
Name: Label, dtype: int64

```



```
In [54]: df_scalar.groupby('Label').mean()
```

```
Out[54]:
```

	Age	Edu	Years Employed	Income	Card Debt	Other Debt	DebtIncomeRatio
Label							
0	0.821946	-0.083259	0.849411	0.635237	-0.192339	-0.180070	-0.593642
1	-0.554105	-0.006161	-0.613107	-0.582946	-0.378133	-0.411896	-0.071906
2	0.560271	0.149212	0.712864	0.944557	1.547956	1.640612	1.156391

from above result We can easily check the centroid values by averaging the features in each cluste

conclusion

I have successfully Studied Unsupervised Machine Learning and pratice and implement k- Means Algorithm

```
In [55]: pred1 = kmean.fit_predict(df_scalar)
score = accuracy_score(df_scalar.Label, pred1)
score
```

```
Out[55]: 1.0
```

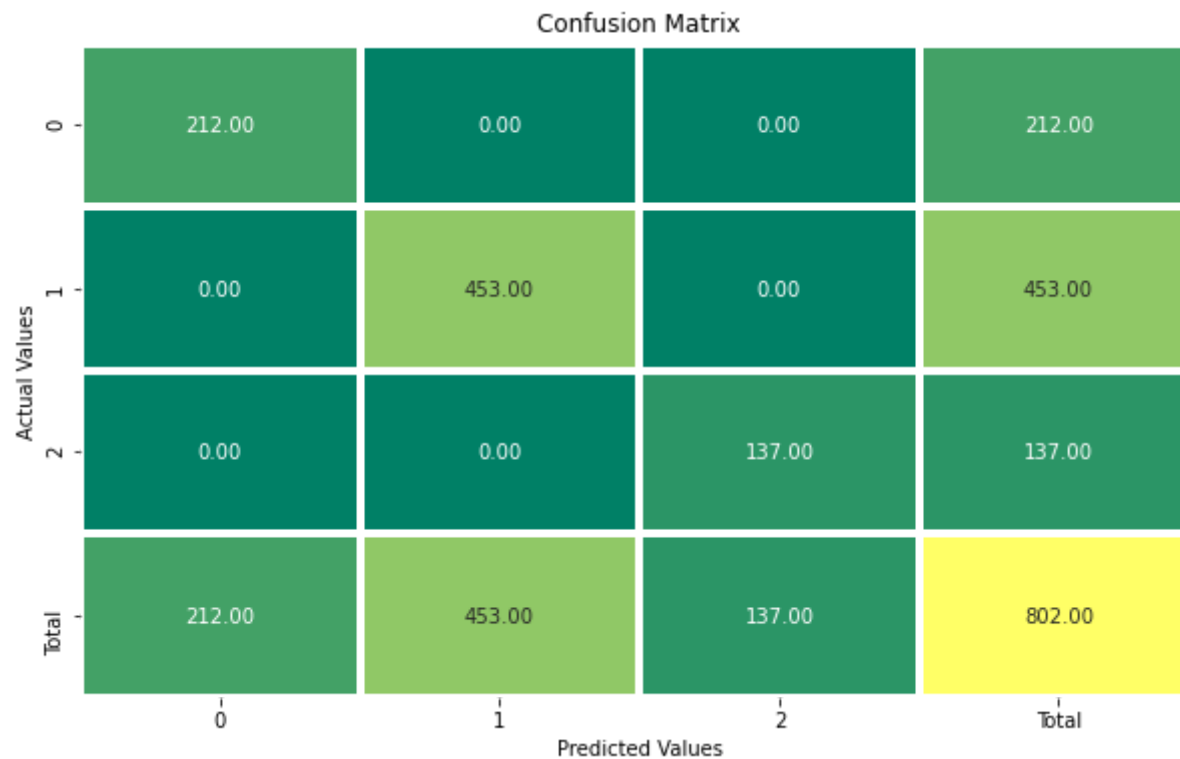
```
In [56]: print(classification_report(df_scalar.Label, pred1))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	212
1	1.00	1.00	1.00	453
2	1.00	1.00	1.00	137
accuracy			1.00	802
macro avg	1.00	1.00	1.00	802
weighted avg	1.00	1.00	1.00	802

```
In [57]: label = [0,1,2]
```

```
In [58]:
```

```
# confusion matrix
cm = confusion_matrix(df_scalar.Label, pred1)
row_sum = cm.sum(axis=0)
cm = np.append(cm, row_sum.reshape(1, -1), axis=0)
col_sum = cm.sum(axis=1)
cm = np.append(cm, col_sum.reshape(-1, 1), axis=1)
labels = label+['Total']
plt.figure(figsize=(10,6))
sns.heatmap(cm, annot=True, cmap='summer', fmt='0.2f', xticklabels=labels,
yticklabels=labels, linewidths=3, cbar=None,)
plt.xlabel('Predicted Values')
plt.ylabel('Actual Values')
plt.title('Confusion Matrix')
plt.show()
```



Conclusion

Thus I have studied unsupervised learning and Successfully Implemented K means algorithm

