

ASSIGNMENT

- Name : Achal Rajesh Mate
- Roll No : 2203541
- Enroll No : MITU20BTCSD001
- Branch : CSE
- Class : TY CSE Is - 3
- Guided By : Prof Nagesh Jadhav Sir

Title:- : Implement CNN for MNIST/CIFAR10 dataset using tensorflow

Objectives:-

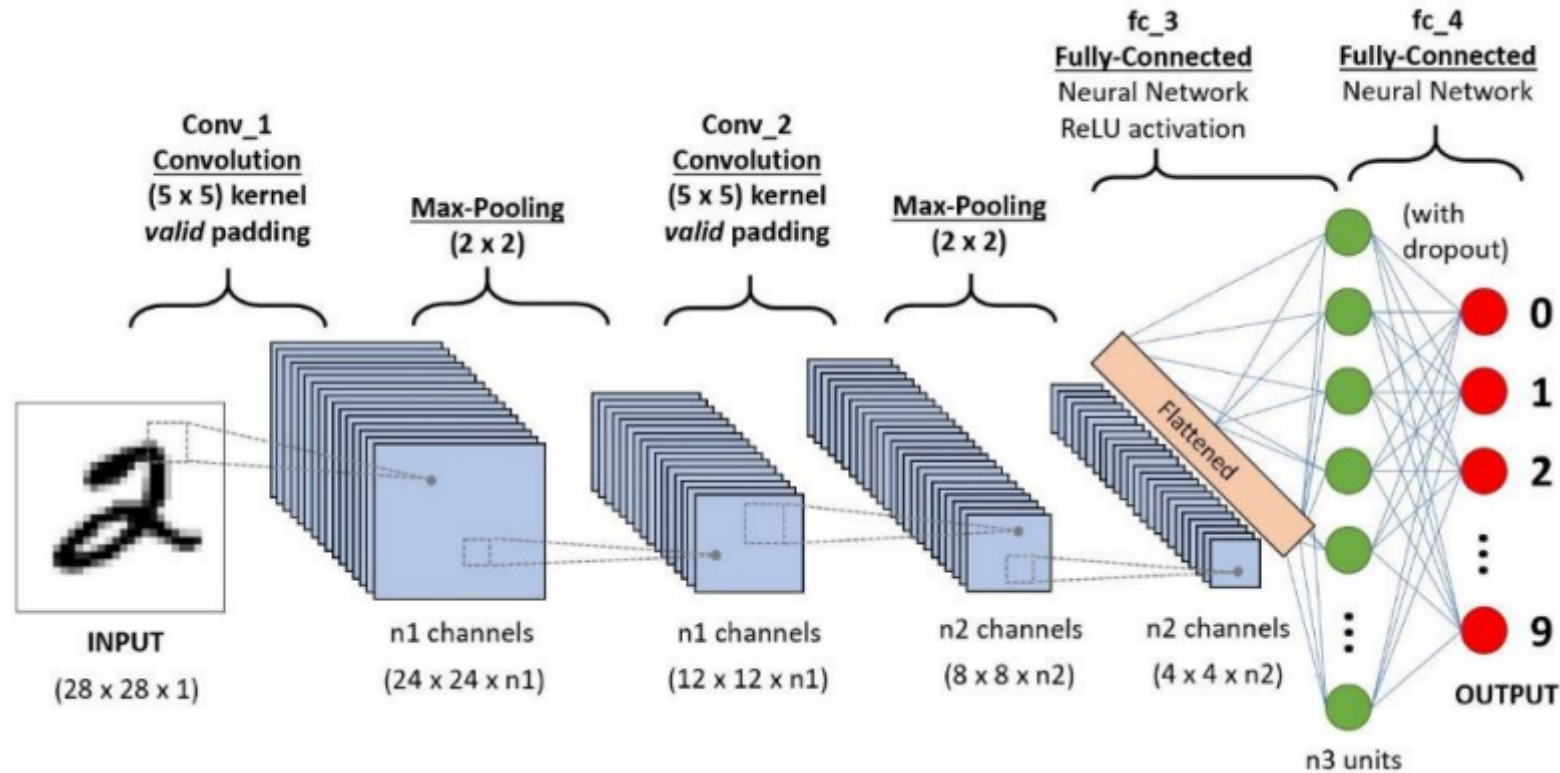
1. To learn basics of deep learning
2. To learn and implement CNN

Theory:

Deep Learning

- Deep learning is an artificial intelligence function that imitates the workings of the human brain in processing data and creating patterns for use in decision making.
- Deep learning is a subset of machine learning in artificial intelligence (AI) that has networks capable of learning unsupervised from data that is unstructured or unlabeled.
- Also known as deep neural learning or deep neural network.
- Deep learning, a subset of machine learning, utilizes a hierarchical level of artificial neural networks to carry out the process of machine learning.
- The artificial neural networks are built like the human brain, with neuron nodes connected together like a web.
- Unlike traditional programs build analysis with data in a linear way, the hierarchical function of deep learning systems enables machines to process data with a nonlinear approach.

CNN Architecture



CNN WORKING

A Convolutional Neural Networks Introduction so to speak.

* Step 1: Convolution Operation :-

The first building block in our plan of attack is convolution operation. In this step, we will touch on feature detectors, which basically serve as the neural network's filters. We will also discuss feature maps, learning the parameters of such maps, how patterns are detected, the layers of detection, and how the findings are mapped out.

* Step 1(b): ReLU Layer :-

The second part of this step will involve the Rectified Linear Unit or ReLU. We will cover ReLU layers and explore how linearity functions in the context of Convolutional Neural Networks. Not necessary for understanding CNN's, but there's no harm in a quick lesson to improve your skills.

* Step 2: Pooling :-

In this part, we'll cover pooling and will get to understand exactly how it generally works. Our nexus here, however, will be a specific type of pooling; max pooling. We'll cover various approaches, though, including mean (or sum) pooling. This part will end with a demonstration made using a visual interactive tool that will definitely sort the whole concept out for you.

* Step 3: Flattening :-

This will be a brief breakdown of the flattening process and how we move from pooled to flattened layers when working with Convolutional Neural Networks.

* Step 4: Full Connection :-

In this part, everything that we covered throughout the section will be merged together. By learning this, you'll get to envision a fuller picture of how Convolutional Neural Networks operate and how the "neurons" that are finally produced learn the classification of images.

Import All Necessary Files

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
```

```
In [27]: import matplotlib.image as mpimg
from tensorflow import keras
from tensorflow.keras import layers
from keras.models import load_model
import tensorflow as tf
import cv2 as cv
from google.colab.patches import cv2_imshow
```

Read the Dataset

```
In [3]: mnist = tf.keras.datasets.mnist
        (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step
11501568/11490434 [=====] - 0s 0us/step


Display Dimensions of training and testing data

```
In [4]: print('X Training shape: ', x_train.shape)
        print('Y Training shape: ', y_train.shape)
        print('X Testing shape: ', x_test.shape)
        print('Y Testing shape: ', y_test.shape)
```

X Training shape: (60000, 28, 28)
Y Training shape: (60000,)
X Testing shape: (10000, 28, 28)
Y Testing shape: (10000,)

```
In [5]: sns.countplot(y_train)
```

Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb770a55410>



```
In [6]: x_train
```

```
Out[6]: array([[0, 0, 0, ..., 0, 0, 0],
               [0, 0, 0, ..., 0, 0, 0],
               [0, 0, 0, ..., 0, 0, 0],
               ...,
               [0, 0, 0, ..., 0, 0, 0],
               [0, 0, 0, ..., 0, 0, 0],
               [0, 0, 0, ..., 0, 0, 0]],

            [[0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             ...,
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0]],

            [[0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             ...,
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0]],

            ...,

            [[0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             ...,
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0]],


            [[0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             ...,
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0]],

            [[0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             ...,
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0],
             [0, 0, 0, ..., 0, 0, 0]]]
```

```
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]], dtype=uint8)
```

Display any one image

```
In [7]: plt.imshow(x_train[200], cmap='gray')
plt.show()
print("Digit in image is ",y_train[200])
```



Digit in image is 1

```
In [8]: input_shape = (28,28,1)
```

View Dimension of image

```
In [9]: x_train[200].shape
```

Out[9]: (28, 28)

```
In [10]: x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255
```

```
In [11]: x_train = np.expand_dims(x_train,-1)
x_test = np.expand_dims(x_test,-1)
```

```
In [12]: batch_size = 128
num_classes = 10
epochs = 5
```

Converts a class vector (integers) to binary class matrix

```
In [13]: y_train = keras.utils.to_categorical(y_train, num_classes) .
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
In [14]: # Layer instances to the constructor:
model = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Dropout(0.5),
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
)		

dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling 2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dropout_1 (Dropout)	(None, 1600)	0
dense (Dense)	(None, 10)	16010

```

=====
Total params: 34,826
Trainable params: 34,826
Non-trainable params: 0

```

```
In [15]: model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
```

```
In [17]: batch_size = 130
num_classes = 10
epochs = 10
```

```
In [18]: history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)
```

```

Epoch 1/10
416/416 [=====] - 68s 163ms/step - loss: 0.0821 - accuracy: 0.9747 - val_loss: 0.0416 - val_accuracy: 0.9887
Epoch 2/10
416/416 [=====] - 64s 153ms/step - loss: 0.0781 - accuracy: 0.9753 - val_loss: 0.0386 - val_accuracy: 0.9883
Epoch 3/10
416/416 [=====] - 49s 118ms/step - loss: 0.0728 - accuracy: 0.9772 - val_loss: 0.0380 - val_accuracy: 0.9897
Epoch 4/10
416/416 [=====] - 46s 111ms/step - loss: 0.0693 - accuracy: 0.9781 - val_loss: 0.0355 - val_accuracy: 0.9908
Epoch 5/10
416/416 [=====] - 45s 109ms/step - loss: 0.0660 - accuracy: 0.9791 - val_loss: 0.0358 - val_accuracy: 0.9902
Epoch 6/10

```



```

416/416 [=====] - 47s 112ms/step - loss: 0.0673 - accuracy: 0.9791 - val_loss: 0.0343 - val_accuracy: 0.9
915
Epoch 7/10
416/416 [=====] - 45s 109ms/step - loss: 0.0616 - accuracy: 0.9805 - val_loss: 0.0317 - val_accuracy: 0.9
912
Epoch 8/10
416/416 [=====] - 47s 112ms/step - loss: 0.0599 - accuracy: 0.9809 - val_loss: 0.0313 - val_accuracy: 0.9
912
Epoch 9/10
416/416 [=====] - 46s 111ms/step - loss: 0.0564 - accuracy: 0.9821 - val_loss: 0.0318 - val_accuracy: 0.9
915
Epoch 10/10
416/416 [=====] - 46s 111ms/step - loss: 0.0565 - accuracy: 0.9820 - val_loss: 0.0285 - val_accuracy: 0.9
920

```

In [19]:

```
print(history.history)
```

```

{'loss': [0.08208952844142914, 0.07807175070047379, 0.07279050350189209, 0.06928857415914536, 0.06604007631540298, 0.0673191398382
1869, 0.06156732887029648, 0.05994338542222977, 0.05640175938606262, 0.056481070816516876], 'accuracy': [0.9746666550636292, 0.975
3147959709167, 0.9771666526794434, 0.9781296253204346, 0.9791296124458313, 0.9790740609169006, 0.9804999828338623, 0.9808518290519
714, 0.9820555448532104, 0.9820370078086853], 'val_loss': [0.04157629981637001, 0.03859942778944969, 0.038046903908252716, 0.03553
558140993118, 0.03579990193247795, 0.03426143527030945, 0.03171133995056152, 0.03130418062210083, 0.03175002709031105, 0.028525330
126285553], 'val_accuracy': [0.9886666536331177, 0.9883333444595337, 0.9896666407585144, 0.9908333420753479, 0.9901666641235352,
0.9915000200271606, 0.9911666512489319, 0.9911666512489319, 0.9915000200271606, 0.9919999837875366]}

```

View Model Performace


In [20]:

```

fig, ax = plt.subplots(2,1)
ax[0].plot(history.history['loss'], color='b', label="Training Loss")
ax[0].plot(history.history['val_loss'], color='r', label="Validation Loss", axes = ax[0])
legend = ax[0].legend(loc='best', shadow=True)

ax[1].plot(history.history['accuracy'], color='b', label="Training Accuracy")
ax[1].plot(history.history['val_accuracy'], color='r', label="Validation Accuracy")
legend = ax[1].legend(loc='best', shadow=True)

```



```
In [21]: test_loss, test_acc = model.evaluate(x_test, y_test)
```


```
313/313 [=====] - 3s 10ms/step - loss: 0.0291 - accuracy: 0.9904
```

```
In [22]: Y_pred = model.predict(x_test)
# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred,axis = 1)
# Convert testing observations to one hot vectors
Y_true = np.argmax(y_test,axis = 1)
# compute the confusion matrix
confusion_mtx = tf.math.confusion_matrix(Y_true, Y_pred_classes)
```

Plot Confusion Matrix

```
In [25]: plt.figure(figsize=(10,10))
sns.heatmap(confusion_mtx, annot=True, fmt='g')
```

```
Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb76b83e090>
```



```
In [31]: import cv2 as cv
from google.colab.patches import cv2_imshow
```

Read new image for testing

```
In [34]: img = cv.imread('five.jfif')
cv2_imshow(img)
```



View Dimension of image

```
In [35]: img.shape
```

```
Out[35]: (233, 216, 3)
```

```
In [37]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

Convert image into black grey

```
In [38]: gray = cv.cvtColor(img, cv.COLOR_BGR2GRAY)
```

```
In [39]: gray.shape
```

```
Out[39]: (233, 216)
```

Change dimension of image

```
In [40]: img_rs = cv.resize(gray, (28, 28))
```

```
In [41]: img_rs.shape
```

```
Out[41]: (28, 28)
```

Display the image

```
In [42]: cv2.imshow(img_rs)
```



```
In [43]: img_rs = np.expand_dims(img_rs,0)
```

```
In [44]: img_rs.shape
```

```
Out[44]: (1, 28, 28)
```

```
In [45]: img_rs = np.expand_dims(img_rs,-1)
```

```
In [46]: img_rs.shape
```

```
Out[46]: (1, 28, 28, 1)
```

make prediction

```
In [47]: num = model.predict(img_rs)
num
```

```
Out[47]: array([[0.96098435, 0.          , 0.          , 0.          , 0.          ,
                0.          , 0.          , 0.          , 0.03901567, 0.          ],
               dtype=float32)
```

```
In [48]: rs = [0,1,2,3,4,5,6,7,8,9]
```

```
In [49]: from numpy.core.fromnumeric import argmax
result = rs[argmax(num)]
```

```
In [50]: result
```

```
Out[50]: 0
```

```
In [51]: model.save('mnist.h5')
```

```
In [52]: model = load_model('mnist.h5')
```

```
In [53]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dropout_1 (Dropout)	(None, 1600)	0

dense (Dense)	(None, 10)	16010
---------------	------------	-------

```
=====
Total params: 34,826
Trainable params: 34,826
Non-trainable params: 0
```

Conclusion

Thus I have learn basics of deep learning and Successfully implement CNN

In []: